

Real-Time Air Quality Monitoring System

Big Data Pipeline with Kafka, Spark, and Multi-Database Architecture

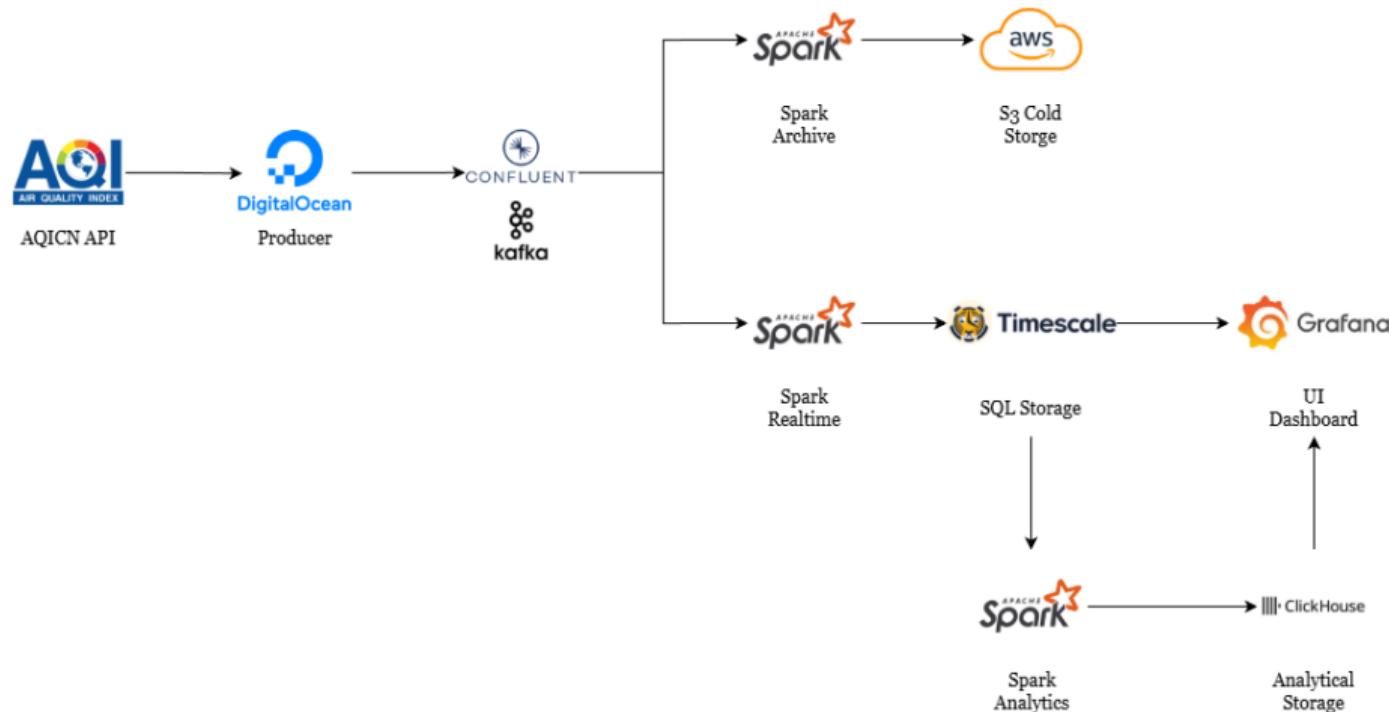
IT4931 - Big Data

December 10, 2025

Project Overview

- **Objective:** Build a scalable real-time air quality monitoring system for 7 Vietnamese cities
- **Data Source:** AQICN API (Air Quality Index Network)
- **Cities:** Hanoi, Ho Chi Minh City, Da Nang, Hai Phong, Can Tho, Nha Trang, Vung Tau
- **Update Frequency:** Every 5 minutes from cloud producers
- **Architecture:** Kappa

System Architecture



Technology Stack

Data Ingestion:

- Kafka (Confluent Cloud)
- Python Producers

Stream Processing:

- Apache Spark 3.5.0
- PySpark Structured Streaming
- Micro-batch (10 min trigger)

Storage:

- TimescaleDB (Time-series)
- ClickHouse (Analytics)
- MinIO (Object Storage)

Visualization:

- Grafana
- PostgreSQL & ClickHouse plugins

Data Pipeline Components

1. Data Producers (DigitalOcean VM)

- 7 independent Python scripts (one per city)
- Fetch data from AQICN API every 5 minutes
- Publish to dedicated Kafka topics: `raw.airquality.<city>`

2. Realtime Processing (Spark Streaming)

- Consumes from Kafka topics
- Enriches data with city metadata (population, region, thresholds)
- Calculates health status and pollution level
- Writes to TimescaleDB hypertables (partitioned by time)

3. Archive Processing (Spark Streaming)

- Archives raw Kafka messages to MinIO
- Parquet format with Snappy compression
- Partitioned by date for compliance and replay

4. Batch Analytics (Spark)

- Hourly aggregations: Reads from TimescaleDB, writes to ClickHouse
- Daily aggregations: Summarizes hourly data
- AQI prediction: ML model using Linear Regression

Database Design - TimescaleDB

Schema: <city>.measurements

- **Purpose:** Store real-time sensor measurements
- **Partitioning:** Hypertable partitioned by processed_time
- **Columns:** AQI, PM2.5, PM10, O3, NO2, temperature, humidity, location
- **Enriched Fields:** health_status, pollution_level, city_metadata

Key Features:

- Automatic time-based partitioning
- Optimized for time-series queries
- PostgreSQL compatibility

Database Design - ClickHouse

Schema: hourly_aggregates, daily_aggregates

- **Purpose:** Fast analytical queries for dashboards
- **Engine:** MergeTree
- **Ordering:** (city, hour_start)

Aggregated Metrics:

- AQI: avg, max, min, stddev
- PM2.5 & PM10: avg, max, min
- Temperature & Humidity: avg
- Record count per hour/day

Deployment Configuration

Cloud Services:

- **Kafka:** Confluent Cloud (`pkc-817wq.ap-east-1.aws.confluent.cloud`)
- **Producers:** DigitalOcean VM (178.128.223.93)

Local Services (Docker):

- TimescaleDB: Port 5432
- ClickHouse: Port 8123 (HTTP), 9003 (Native)
- MinIO: 4-node cluster, Port 9000 (API), 9001 (Console)
- Grafana: Port 3000

Processing Jobs:

- 7 realtime streaming jobs (one per city)
- 7 archive streaming jobs
- 7 hourly analytics jobs

Data Flow Example

Realtime Processing Pipeline:

- ① Producer fetches AQI data from AQICN API (every 5 min)
- ② Data published to Kafka topic: `raw.airquality.hanoi`
- ③ Spark Streaming consumes JSON messages
- ④ Parse and validate data schema
- ⑤ Join with city metadata (broadcast)
- ⑥ Calculate derived fields:
 - `health_status`: Good/Moderate/Unhealthy/Hazardous
 - `pollution_level`: Low/Medium/High/Severe
- ⑦ Write to TimescaleDB: `hanoi_measurements`
- ⑧ Checkpoint state every 10 minutes

Code Structure

Source Code:

- src/common/config.py
- src/ingestion/producer.py
- src/processing/realtme.py
- src/processing/archive.py
- src/processing/analytics_hourly.py
- src/processing/analytics_daily.py
- src/processing/aqi_prediction.py

Run Scripts (per city):

- scripts/produce_<city>.py
- scripts/run_realtime_<city>.py
- scripts/run_archive_<city>.py
- scripts/run_analytics_hourly_<city>.py
- scripts/run_aqi_prediction_<city>.py

Configuration:

- docker-compose.yaml: Infrastructure services
- .env: Credentials and endpoints
- data/city_metadata.json: Static reference data

Current Status

Completed:

- ✓ Kafka producers (7 cities)
- ✓ Docker infrastructure
- ✓ Realtime streaming pipeline
- ✓ TimescaleDB integration
- ✓ City metadata management
- ✓ Per-city run scripts

In Progress:

- Archive processing (MinIO)
- Hourly analytics aggregation
- ClickHouse integration
- Grafana dashboards
- AQI prediction models

Planned Visualizations (Grafana)

5 Dashboard Views:

① Real-time City Dashboard

- Current AQI, PM2.5, PM10 levels
- Health status indicators
- 24-hour trend charts

② Historical Trends

- Weekly/monthly AQI patterns
- Pollution level comparisons

③ City Comparison

- Side-by-side metrics for all 7 cities
- Regional analysis (North/Central/South)

Planned Visualizations

④ AQI Predictions

- ML-based forecasts (next 24 hours)
- Confidence intervals
- Feature importance

⑤ Summary Statistics

- Daily/weekly aggregates
- Best/worst air quality rankings
- Health risk alerts

Future Enhancements

Short-term:

- Complete Grafana dashboard implementation
- Deploy ML prediction models
- Set up automated alerting

Long-term:

- Add more cities and sensors
- Implement Delta Lake for data lake
- Build REST API for external access
- Mobile app for public access
- Integrate weather data for correlation analysis

Conclusion

Project Achievements:

- Built end-to-end real-time data pipeline
- Integrated 5 different technologies (Kafka, Spark, TimescaleDB, ClickHouse, MinIO)
- Implemented Lambda architecture
- Handled 7 cities with scalable per-city processing
- Cloud-native design with local development environment

Technical Skills Demonstrated:

- Distributed stream processing
- Time-series database optimization
- Docker containerization
- Cloud service integration
- Data pipeline orchestration

Thank You!